

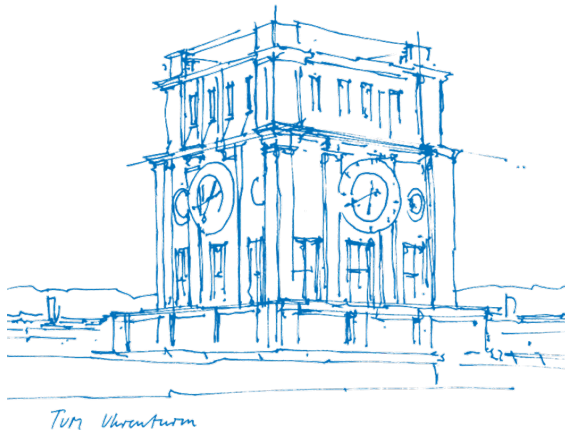
Docker resource management

ACA Student Presentation

Roberto Castellotti

TUM School of Computation, Information and
Technology
Technical University of Munich

December 14th, 2022



Quick questions

- Who has ever heard of Docker?
- Who used Docker at least once?
- Who uses Docker in their daily life?
- Who knows that tools similar to Docker exist? (or focusing on subtasks)
- Who has seen this slide at least 5 times?

Containers vs VMs (once again)




Figure 1 from: <https://www.docker.com/resources/what-container/>

"Docker"

The Open Container Initiative is an open governance structure for the express purpose of creating open industry standards around container formats and runtimes.



podman



buildah



skopeo



HashiCorp

Packer



Kaniko

containerd



Finch

How do containers provide isolation?

"Applications are safer in containers and Docker provides the strongest default isolation capabilities in the industry" ¹

- linux namespaces (limits what you see)
- cgroups (limits what you can use use)

¹<https://www.docker.com/resources/what-container/>

Namespaces are a feature of the Linux kernel that partitions kernel resources such that one set of processes sees one set of resources while another set of processes sees a different set of resources.... Examples of such resources are process IDs, hostnames, user IDs, file names, and some names associated with network access, and interprocess communication.²

- inspired by the wider namespace functionality used heavily throughout Plan 9 from Bell Labs.
- 2002: initial work on namespaces introduced in Linux mainline kernel 2.4.19
- February 2013: user namespaces are introduced in Linux mainline kernel 3.8 (we now have adequate container support functionality)

²https://en.wikipedia.org/wiki/Linux_namespaces

Type of Linux Namespaces

- **User** namespaces isolate User and group IDs
- **PID** namespaces isolate process IDs (nested trees; PID 1 in a subtree)
- **Mount** namespaces isolate mount points
- **Network** namespaces isolate Network devices, stacks, ports
- some others, check man page

Let's get our hands dirty, user namespace(1/2)

```
# outside the namespace
rc@s369 ~/t/aca> whoami
rc
rc@s369 ~/t/aca> id
uid=1000(rc) gid=1000(rc) groups=1000(rc), ...
rc@s369 ~/t/aca> sudo unshare --user

# inside the namespace
nobody@s369:/home/rc/tum-exams/aca$ whoami
nobody
nobody@s369:/home/rc/tum-exams/aca$ id
uid=65534(nobody) gid=65534(nogroup) groups=65534(nogroup)
nobody@s369:/home/rc/tum-exams/aca$
```

Who is nobody? "If a user ID has no mapping inside the namespace, then system calls that return user IDs return the value defined in the file `/proc/sys/kernel/overflowuid`, which on a standard system defaults to the value 65534."³

³<https://lwn.net/Articles/532593/>

Let's get our hands dirty, user namespace (2/2)

```
# outside the namespace
rc@s369 ~/t/aca> whoami
rc
rc@s369 ~→ id
uid=1000(rc) gid=1000(rc) groups=1000(rc), ...
rc@s369 ~→ sudo unshare --user --map-root-user
# inside the namespace
root@s369:/home/rc# whoami
root
root@s369:/home/rc# id
uid=0(root) gid=0(root) groups=0(root)
```

This is probably what you need if you are creating a user namespace.

Let's get our hands dirty, mount namespace

```
rc@s369 ~→ sudo unshare --mount
```

```
# inside the namespace
```

```
root@s369:/home/rc# mount -t tmpfs tmpfs /mnt
```

```
root@s369:/home/rc# echo "I love ACA student presentations!" > /mnt/aca.txt
```

```
root@s369:/home/rc# cat /mnt/aca.txt
```

```
I love ACA student presentations!
```

```
root@s369:/home/rc#
```

```
logout
```

```
# outside of namespace
```

```
rc@s369 ~→ ls -l /mnt/aca.txt
```

```
ls: cannot access '/mnt/aca.txt': No such file or directory
```

```
rc@s369 ~ [2]>
```

Let's get our hands dirty, net namespace (1/2)

```
# outside the namespace
```

```
rc@s369 ~→ ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000  
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
    inet 127.0.0.1/8 scope host lo
```

```
        valid_lft forever preferred_lft forever
```

```
    inet6 ::1/128 scope host
```

```
        valid_lft forever preferred_lft forever
```

```
2: enp0s31f6: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc fq_codel state DOWN group default
```

```
    link/ether 54:e1:ad:5d:1c:6e brd ff:ff:ff:ff:ff:ff
```

```
.... i should probably clean this up :)
```

```
rc@s369 ~→ ip route
```

```
default via 131.159.223.254 dev wlp4s0 proto dhcp metric 600
```

```
131.159.192.0/19 dev wlp4s0 proto kernel scope link src 131.159.218.200 metric 600
```

```
169.254.0.0/16 dev br-0d5fb61239aa scope link metric 1000
```

```
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1 linkdown
```

```
172.18.0.0/16 dev br-0d5fb61239aa proto kernel scope link src 172.18.0.1
```

Let's get our hands dirty, net namespace (2/2)

```
rc@s369 ~→ sudo unshare --net /bin/bash
# inside the namespace
root@s369:/home/rc# ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
root@s369:/home/rc# iptables --list-rules
-P INPUT ACCEPT
-P FORWARD ACCEPT
-P OUTPUT ACCEPT
root@s369:/home/rc# ip route
root@s369:/home/rc#
```

Let's get our hands dirty, pid namespace (1/2)

```
rc@s369 ~→ ps -aux | head -2
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	168412	13844	?	Ss	09:56	0:04	/sbin/init splash

```
rc@s369 ~→ ps -aux | wc -l
```

```
324
```

```
rc@s369 ~→ sudo unshare --fork --pid /bin/bash
```

```
root@s369:/home/rc# ps -aux | wc -l
```

```
328
```

```
root@s369:/home/rc# ps -aux | head -2
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	168412	13844	?	Ss	09:56	0:04	/sbin/init splash

```
rc@s369 ~→ sudo unshare --fork --pid --mount-proc /bin/bash
```

```
root@s369:/home/rc# ps -aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	10236	4132	pts/4	S	18:53	0:00	/bin/bash
root	8	0.0	0.0	12940	3720	pts/4	R+	18:53	0:00	ps -aux

Why this happens?

Let's get our hands dirty, pid namespace (2/2)

The ps program uses the procfs virtual file system to obtain information about the current processes in the system. This filesystem is mounted in the /proc directory. However, in the new namespace this mountpoint describes the processes from the root PID namespace.

```
rc@cs369:~$ echo $$
```

```
45460
```

```
rc@cs369:~$ ls /proc/45460
```

arch_status	cwd	mem	patch_state	stat
attr	environ	mountinfo	personality	statm
autogroup	exe	mounts	projid_map	status
auxv	fd	mountstats	root	syscall
cgroup	fdinfo	net	sched	task
clear_refs	gid_map	ns	schedstat	timens_offsets
cmdline	io	numa_maps	sessionid	timers
comm	limits	oom_adj	setgroups	timerslack_ns
coredump_filter	loginuid	oom_score	smaps	uid_map
cpu_resctrl_groups	map_files	oom_score_adj	smaps_rollup	wchan
cpuset	maps	pagemap	stack	

```
rc@cs369:~$ ls /proc/45460/ns
```

cgroup	mnt	pid	time	user
ipc	net	pid_for_children	time_for_children	uts

Cgroups is a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes.⁴

- 2006: engineers at Google start working of this feature
- January 2008: cgroups functionality is merged in the Linux mainline kernel (2.6.24)
- December 2016: cgroups v2 is merged in Linux mainline kernel (4.5) with some improvements

⁴<https://en.wikipedia.org/wiki/Cgroups>

Cgroup controllers

- **cpu**: grants a minimum number of "CPU shares" when system is busy
 - from 3.2: CPU "bandwidth" control
- **devices**: supports controlling which processes may create (mknod) devices as well as open them for reading or writing.
- **cpuacct**: provides accounting for CPU usage by groups of processes.
- **pids**: permits limiting the number of process that may be created in a cgroup (and its descendants).
- **net_cls**: places a classid on network packets created by a cgroup.
- **memory**: supports reporting and limiting of process memory, kernel memory, and swap used by cgroups.
- **blkio**: controls and limits access to specified block devices.
- **hugetlb, rdma, cpuset, cpu_prio, net_prio, perf_event, freezer**

Memory controller hands on

```
root@s369:/# mkdir /sys/fs/cgroup/memory/cgtest
root@s369:/# echo 5000 > /sys/fs/cgroup/memory/cgtest/memory.limit_in_bytes
root@s369:/# cat /sys/fs/cgroup/memory/cgtest/memory.limit_in_bytes
4096
```

now let's create a simple script test.sh

```
#!/bin/bash
```

```
while [ 1 ]; do
    echo "I love Advanced Computer Networking Student Presentations!"
    sleep 20
done
```

Memory controller hands on

```
rc@s369:~$ sh ~/test.sh &
```

```
[1] 9668
```

```
root@s369:/# echo 9668 > /sys/fs/cgroup/memory/cgtest/cgroup.procs
```

```
rc@s369 ~$ ps 9668
```

PID	TTY	STAT	TIME	COMMAND
9668	pts/2	S	0:00	sh /home/rc/test.sh

```
rc@s369 ~$ ps -o cgroup 9668
```

```
CGROUP
```

```
13:memory:/cgtest,11:pids:/user.slice/user-1000.slice/user@1000.service...
```

after some time... test.sh

```
rc@s369 ~$ ps 9668
```

PID	TTY	STAT	TIME	COMMAND
-----	-----	------	------	---------

What happened?

Memory controller hands on

```
rc@s369:~$ sh ~/test.sh &
```

```
[1] 9668
```

```
root@s369:/# echo 9668 > /sys/fs/cgroup/memory/cgtest/cgroup.procs
```

```
rc@s369 ~$ ps 9668
```

PID	TTY	STAT	TIME	COMMAND
9668	pts/2	S	0:00	sh /home/rc/test.sh

```
rc@s369 ~$ ps -o cgroup 9668
```

```
CGROUP
```


```
13:memory:/cgtest,11:pids:/user.slice/user-1000.slice/user@1000.service...
```

after some time... test.sh

```
rc@s369 ~$ ps 9668
```

PID	TTY	STAT	TIME	COMMAND
-----	-----	------	------	---------

What happened? There is a new sheriff in town! oom-killer!



further readings/talks

- <https://github.com/containerd/containerd/blob/main/docs/getting-started.md>
- <https://lwn.net/Articles/531114/> (the entire series)
- <https://blog.quarkslab.com/digging-into-linux-namespaces-part-1.html>
- `man {namespaces,cgroups}`
- <https://drewdevault.com/2022/11/12/In-praise-of-Plan-9.html>
- @jpetazzo - Cgroups, namespaces, and beyond: what are containers made from?
- Containers unplugged: Linux namespaces - Michael Kerrisk
- <https://docs.kernel.org/admin-guide/cgroup-v1/cgroups.html>